

Topic_4_Basic_programming_concepts

March 4, 2025

0.1 Data Series

1. In Pandas, series are labeled/indexed one-dimensional arrays.
2. The elements of the series could be integer, float, string, objects, etc.
3. The index may not be unique.
4. A collection of Data Series is called as Data Frame.
5. Our focus in this chapter will be on Data Frames.

1 DataFrame (DF)

Collection of series, where each column is a series.

1.1 Creating DataFrame

In order to create a data from, following things are typically needed:

1. Data in n-dimensional array format.
2. Optional column title/header

An optional row index can be given too. In addition to the above, there are many parameters, which can be read from pydata docs.

1.1.1 Create DF with Given n-D Array

Data:

$$\begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} \quad (1)$$

col_names: ['col 1', 'col 2']

row_index: ['row 1', 'row 2']

```
[9]: %pip install pandas
```

```
Requirement already satisfied: pandas in  
/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (2.2.3)  
Requirement already satisfied: numpy>=1.23.2 in  
/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (from  
pandas) (2.2.3)  
Requirement already satisfied: python-dateutil>=2.8.2 in  
/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (from
```

pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in
/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (from
pandas) (2025.1)
Requirement already satisfied: tzdata>=2022.7 in
/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (from
pandas) (2025.1)
Requirement already satisfied: six>=1.5 in
/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (from
python-dateutil>=2.8.2->pandas) (1.17.0)

[notice] A new release of pip is
available: 24.0 -> 25.0.1
[notice] To update, run:
pip install --upgrade pip
Note: you may need to restart the kernel to use updated packages.

```
[10]: import pandas as pd
df = pd.DataFrame([[ 'a11', 'a12'], [ 'a21', 'a22']],
                  index=[ 'row 1', 'row 2'],
                  columns=[ 'Course', 'Major'])

display(df)
```

	Course	Major
row 1	a11	a12
row 2	a21	a22

1.1.2 Create DF with Random n-D Array

Data: random 3 columns 1000 rows numbers between 0 and 1.
col_names: ['col 1', 'col 2', 'col 3']

```
[11]: import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.rand(1000,3),
                  columns=[ 'col 1', 'col 2', 'col 3'])

display(df)
```

	col 1	col 2	col 3
0	0.882961	0.598013	0.231091
1	0.476830	0.598566	0.082018
2	0.524957	0.443243	0.648673
3	0.849683	0.824520	0.425529
4	0.868353	0.729919	0.060088
..
995	0.737923	0.805233	0.051131

```

996  0.381251  0.677263  0.458775
997  0.056614  0.426475  0.718898
998  0.398317  0.911468  0.236637
999  0.462512  0.609815  0.950087

```

[1000 rows x 3 columns]

1.1.3 Create DF from CSV

Data: Basic-1.csv

col_names: Use from the first row

1. Display the above data.
2. Obtain the summary statistics:
 - For numeric data (count, mean, std, min, max)
 - For non-numeric data (count, unique-values, most occurring and its frequency)

```

[12]: #1. Display the above data.
import pandas as pd
df = pd.read_csv('data/Basic-1.csv', delimiter=',')#the path to the data file.

## Path relative to the .ipynb file
# df = pd.read_csv('Basic-101.csv',
    ↪delimiter=',', header=None, names=['Gender', 'Job Type', "Province"])

# display data
display(df)

#2. Obtain the summary statistics:
# summary statistics for non-numeric columns is available by default only when
    ↪all columns are non-numeric
display(df.describe())

```

	Gender	Job Type	Province
0	M	Pink-collar	Hejaz
1	M	White-collar	Central
2	M	Pink-collar	Hejaz
3	M	Pink-collar	Hejaz
4	M	Gold-collar	Eastern
...
1273	M	Blue-collar	Central
1274	M	Pink-collar	Hejaz
1275	F	Blue-collar	Hejaz
1276	F	White-collar	Central
1277	F	Blue-collar	Eastern

[1278 rows x 3 columns]

Gender	Job Type	Province
--------	----------	----------

count	1278	1278	1278
unique	2	4	3
top	M	White-collar	Central
freq	875	568	531

1.1.4 Create DF from XLSX

Data: Basic-2.xlsx

col_names: There are 17 columns, defined respectively as: 1. dur: duration of agreement 2. wage1.wage : wage increase in first year of contract 3. wage2.wage : wage increase in second year of contract 4. wage3.wage : wage increase in third year of contract 5. cola : cost of living allowance 6. hours.hrs : number of working hours during week 7. pension : employer contributions to pension plan 8. stby_pay : standby pay 9. shift_diff : shift differential : supplement for work on II and III shift 10. educ_allw.boolean : education allowance 11. holidays : number of statutory holidays 12. vacation : number of paid vacation days 13. lngtrm_disabil.boolean : employer's help during employee longterm disability 14. dntl_ins : employers contribution towards the dental plan 15. bereavement.boolean : employer's financial contribution towards the covering the costs of bereavement 16. empl_hplan : employer's contribution towards the health plan 17. empl_hplan : employer's contribution towards the health plan

For the above data, do the following: 1. Display the above data. 2. Obtain the summary statistics:
- For numeric data (count, mean, std, min, max) - For non-numeric data (count, unique-values, most occurring and its frequency)

```
[13]: %pip install openpyxl
```

```
Requirement already satisfied: openpyxl in
/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (3.1.5)
Requirement already satisfied: et-xmlfile in
/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (from
openpyxl) (2.0.0)
```

[notice] A new release of pip is available: 24.0 -> 25.0.1

[notice] To update, run:

```
pip install --upgrade pip
```

Note: you may need to restart the kernel to use updated packages.

```
[14]: # 1. Display the above data.
```

```
import pandas as pd
```

```
df = pd.read_excel('data/Basic-2.xlsx')
```

```
#display(df.head())
```

```
display(df.tail())
```

```
#df.sample(5)
```

	c1	c2	c3	c4	c5	c6	c7	c8	c9	c10	c11	\
35	2.0	2.0	2.0	NaN	none	40.0	none	NaN	NaN	no	11.0	
36	1.0	2.0	NaN	NaN	tc	40.0	ret_allw	4.0	0.0	no	11.0	

37	1.0	2.8	NaN	NaN	none	38.0	empl_contr	2.0	3.0	no	9.0
38	3.0	2.0	2.5	2.0	NaN	37.0	empl_contr	NaN	NaN	NaN	10.0
39	2.0	4.5	4.0	NaN	none	40.0		NaN	NaN	4.0	12.0

		c12	c13	c14	c15	c16	c17
35		average	yes	none	yes	full	bad
36		generous	no	none	no	none	bad
37	below	average	yes	half	NaN	none	bad
38		average	NaN	NaN	yes	none	bad
39		average	yes	full	yes	half	good

```
[15]: # 2. Obtain the summary statistics:

#display(df.describe())
# by default the describe gives summary of the numeric columns, when there are
↳ mixed columns
#df.describe(include='all') # can be used, but too many NaNs
```

```
[16]: # to get summary of non-numeric columns

# Identify non-numeric columns
cat_columns= df.select_dtypes(include='object').columns

# Show summary of non-numeric columns
df[cat_columns].describe()

## or in one step
# df.select_dtypes(include='object').describe()
```

```
[16]:
```

	c5	c7	c10		c12	c13	c14	c15	c16	c17
count	24	18	18		37	16	25	20	24	40
unique	3	3	2		3	2	3	2	3	2
top	none	none	no	below average	yes	half	yes	full	good	
freq	14	8	11		14	11	11	18	12	26

1.2 Accessing Elements

`.loc` and `.iloc` methods can be used for accessing the elements of a DF.

1.2.1 Accessing via `.loc`

Consider the dataset given in Basic-1.csv. Display second row, second column element.

```
[17]: import pandas as pd
df = pd.read_csv('data/Basic-1.csv', delimiter=',')
print(df.loc[1, 'Job Type'])
```

White-collar

1.2.2 Accessing via .iloc

Consider the dataset given in Basic-1.csv. Display second row, second column element.

```
[18]: import pandas as pd
df = pd.read_csv('data/Basic-1.csv', delimiter=',')
display(df.head())
print(df.iloc[1,1])
print(df.loc[1,'Job Type'])
```

	Gender	Job Type	Province
0	M	Pink-collar	Hejaz
1	M	White-collar	Central
2	M	Pink-collar	Hejaz
3	M	Pink-collar	Hejaz
4	M	Gold-collar	Eastern

White-collar

White-collar

1.3 DF Slicing

1. Slicing is nothing but extracting a part of the DF.
2. The part could be contiguous or broken chunks of the DF.

1.3.1 Single Column Slice

A single column can be extracted using the following methods: 1. df.COL_NAME 2. df[COL_NAME]

Consider the dataset given in Basic-2.xlsx. Display first column.

```
[19]: import pandas as pd

df = pd.read_excel('data/Basic-2.xlsx')
display(df.head())
display(df['c1'])

## Other styles to select a column
# display(df.iloc[:,0])
# display(df.loc[:, 'c1'])

## Following code can be used to create column names when dataset does not
    contain any column names.
# col_name=[''.join(['c',str(i+1)]) for i in range(17)]
## Following options can be used to add column names to the dataset
# header=None, names=col_name
```

	c1	c2	c3	c4	c5	c6	c7	c8	c9	c10	c11	\
0	1.0	5.0	NaN	NaN	NaN	40.0	NaN	NaN	2.0	NaN	11.0	
1	2.0	4.5	5.8	NaN	NaN	35.0	ret_allw	NaN	NaN	yes	11.0	

2	NaN	NaN	NaN	NaN	NaN	38.0	empl_contr	NaN	5.0	NaN	11.0
3	3.0	3.7	4.0	5.0	tc	NaN		NaN	NaN	NaN	yes
4	3.0	4.5	4.5	5.0	NaN	40.0		NaN	NaN	NaN	NaN

		c12	c13	c14	c15	c16	c17
0		average	NaN	NaN	yes	NaN	good
1	below	average	NaN	full	NaN	full	good
2		generous	yes	half	yes	half	good
3		NaN	NaN	NaN	yes	NaN	good
4		average	NaN	half	yes	half	good

0	1.0
1	2.0
2	NaN
3	3.0
4	3.0
5	2.0
6	3.0
7	3.0
8	2.0
9	1.0
10	3.0
11	2.0
12	2.0
13	3.0
14	1.0
15	2.0
16	1.0
17	1.0
18	1.0
19	2.0
20	2.0
21	2.0
22	3.0
23	2.0
24	1.0
25	3.0
26	2.0
27	2.0
28	2.0
29	3.0
30	3.0
31	3.0
32	2.0
33	2.0
34	3.0
35	2.0
36	1.0

```

37     1.0
38     3.0
39     2.0
Name: c1, dtype: float64

```

1.3.2 Slicing via .loc

Consider the dataset given in Basic-2.xlsx. Display first 10 rows from the second, third, fourth and fifth column.

```

[20]: import pandas as pd
df = pd.read_excel('data/Basic-2.xlsx')

display(df.head())

display(df.loc[0:9,'c2':'c5']) # ends are inclusive

# display(df.loc[0:9:2,'c2':'c5':2]) # ends are inclusive and steps can be
↳used too

```

	c1	c2	c3	c4	c5	c6	c7	c8	c9	c10	c11	\
0	1.0	5.0	NaN	NaN	NaN	40.0	NaN	NaN	2.0	NaN	11.0	
1	2.0	4.5	5.8	NaN	NaN	35.0	ret_allw	NaN	NaN	yes	11.0	
2	NaN	NaN	NaN	NaN	NaN	38.0	empl_contr	NaN	5.0	NaN	11.0	
3	3.0	3.7	4.0	5.0	tc	NaN	NaN	NaN	NaN	yes	NaN	
4	3.0	4.5	4.5	5.0	NaN	40.0	NaN	NaN	NaN	NaN	12.0	

	c12	c13	c14	c15	c16	c17
0	average	NaN	NaN	yes	NaN	good
1	below average	NaN	full	NaN	full	good
2	generous	yes	half	yes	half	good
3	NaN	NaN	NaN	yes	NaN	good
4	average	NaN	half	yes	half	good

	c2	c3	c4	c5
0	5.0	NaN	NaN	NaN
1	4.5	5.8	NaN	NaN
2	NaN	NaN	NaN	NaN
3	3.7	4.0	5.0	tc
4	4.5	4.5	5.0	NaN
5	2.0	2.5	NaN	NaN
6	4.0	5.0	5.0	tc
7	6.9	4.8	2.3	NaN
8	3.0	7.0	NaN	NaN
9	5.7	NaN	NaN	none

```

[21]: #pip install openpyxl

```


1.3.3 Slicing via .iloc

Consider the dataset given in Basic-2.xlsx. Display first 10 rows from the second, third, fourth and fifth column.

```
[22]: import pandas as pd
df = pd.read_excel('data/Basic-2.xlsx')

display(df.iloc[:10,1:5])
```

	c2	c3	c4	c5
0	5.0	NaN	NaN	NaN
1	4.5	5.8	NaN	NaN
2	NaN	NaN	NaN	NaN
3	3.7	4.0	5.0	tc
4	4.5	4.5	5.0	NaN
5	2.0	2.5	NaN	NaN
6	4.0	5.0	5.0	tc
7	6.9	4.8	2.3	NaN
8	3.0	7.0	NaN	NaN
9	5.7	NaN	NaN	none

```
[23]: ## Example of loc and iloc
df = pd.DataFrame([[ 'a11', 'a12'], [ 'a21', 'a22']],
                  index=[ 'row 1', 'row 2'],
                  columns=[ 'col 1', 'col-2'])

display(df)

# display(df.loc[:, 'col 1']) # Select one column

# display(df['col 1']) # Select one column

# display(df.col-2) # # Select one column, problem is Spaces or Special
↳ characters

# display(df.loc['row 2', 'col 1'])

# display(df.loc['row 2', :])

# display(df['row 2']) ## check...

# display(df.iloc[:, 0]) # Select one column

# display(df.iloc[1, 0]) # row 2 column 1
```

col 1 col-2

```
row 1    a11    a12
row 2    a21    a22
```

1.4 DF Sorting

1. DataFrames can be sorted w.r.t values in the columns.
2. One or more columns can be used for sorting.

Consider the dataset given in Basic-2.xlsx. Sort the rows according to the first column, in ascending order. Repeat the sort in descending order.

```
[24]: import pandas as pd

#Read data
df = pd.read_excel('data/Basic-2.xlsx')
display(df.head())

# # sort by single column in ascending order
df.sort_values(by=['c1'], inplace=True)
display(df.head())

# # sort by single column in descending order
df.sort_values(by=['c1'], inplace=True, ascending=False)
display(df.head())
```

	c1	c2	c3	c4	c5	c6	c7	c8	c9	c10	c11	\
0	1.0	5.0	NaN	NaN	NaN	40.0	NaN	NaN	2.0	NaN	11.0	
1	2.0	4.5	5.8	NaN	NaN	35.0	ret_allw	NaN	NaN	yes	11.0	
2	NaN	NaN	NaN	NaN	NaN	38.0	empl_contr	NaN	5.0	NaN	11.0	
3	3.0	3.7	4.0	5.0	tc	NaN	NaN	NaN	NaN	yes	NaN	
4	3.0	4.5	4.5	5.0	NaN	40.0	NaN	NaN	NaN	NaN	12.0	

	c12	c13	c14	c15	c16	c17
0	average	NaN	NaN	yes	NaN	good
1	below average	NaN	full	NaN	full	good
2	generous	yes	half	yes	half	good
3	NaN	NaN	NaN	yes	NaN	good
4	average	NaN	half	yes	half	good

	c1	c2	c3	c4	c5	c6	c7	c8	c9	c10	c11	\
0	1.0	5.0	NaN	NaN	NaN	40.0	NaN	NaN	2.0	NaN	11.0	
16	1.0	2.8	NaN	NaN	NaN	35.0	NaN	NaN	2.0	NaN	12.0	
14	1.0	3.0	NaN	NaN	none	36.0	NaN	NaN	10.0	no	11.0	
9	1.0	5.7	NaN	NaN	none	40.0	empl_contr	NaN	4.0	NaN	11.0	
18	1.0	2.0	NaN	NaN	none	38.0	none	NaN	NaN	yes	11.0	

	c12	c13	c14	c15	c16	c17
0	average	NaN	NaN	yes	NaN	good
16	below average	NaN	NaN	NaN	NaN	good

14		generous	NaN	NaN	NaN	NaN	good
9		generous	yes	full	NaN	NaN	good
18		average	no	none	no	none	bad

	c1	c2	c3	c4	c5	c6	c7	c8	c9	c10	c11	\
22	3.0	3.5	4.0	4.6	tcf	27.0	NaN	NaN	NaN	NaN	NaN	
34	3.0	2.0	2.5	2.1	tc	40.0	none	2.0	1.0	no	10.0	
29	3.0	2.0	2.5	NaN	NaN	35.0	none	NaN	NaN	NaN	10.0	
38	3.0	2.0	2.5	2.0	NaN	37.0	empl_contr	NaN	NaN	NaN	10.0	
30	3.0	4.5	4.5	5.0	none	40.0	NaN	NaN	NaN	no	11.0	

		c12	c13	c14	c15	c16	c17
22		NaN	NaN	NaN	NaN	NaN	good
34	below	average	no	half	yes	full	bad
29		average	NaN	NaN	yes	full	bad
38		average	NaN	NaN	yes	none	bad
30		average	NaN	half	NaN	NaN	good

Consider the dataset given in Basic-2.xlsx. Sort the rows according to the first and then sixth column, where ascending in the first and descending in the sixth column. Display only first 6 columns.

```
[25]: # sort by multiple columns

df.sort_values(by=['c1','c6'], inplace=True, ascending=[True, False])

display(df.iloc[:,0:6])
```

	c1	c2	c3	c4	c5	c6
0	1.0	5.0	NaN	NaN	NaN	40.0
17	1.0	2.1	NaN	NaN	tc	40.0
9	1.0	5.7	NaN	NaN	none	40.0
36	1.0	2.0	NaN	NaN	tc	40.0
24	1.0	6.0	NaN	NaN	NaN	38.0
18	1.0	2.0	NaN	NaN	none	38.0
37	1.0	2.8	NaN	NaN	none	38.0
14	1.0	3.0	NaN	NaN	none	36.0
16	1.0	2.8	NaN	NaN	NaN	35.0
23	2.0	4.5	4.0	NaN	NaN	40.0
35	2.0	2.0	2.0	NaN	none	40.0
12	2.0	3.5	4.0	NaN	none	40.0
33	2.0	4.0	5.0	NaN	none	40.0
21	2.0	2.5	3.0	NaN	NaN	40.0
39	2.0	4.5	4.0	NaN	none	40.0
11	2.0	6.4	6.4	NaN	NaN	38.0
8	2.0	3.0	7.0	NaN	NaN	38.0
32	2.0	2.5	2.5	NaN	NaN	38.0
20	2.0	4.3	4.4	NaN	NaN	38.0
28	2.0	5.0	4.0	NaN	none	37.0

15	2.0	4.5	4.0	NaN	none	37.0
19	2.0	4.0	5.0	NaN	tcf	35.0
1	2.0	4.5	5.8	NaN	NaN	35.0
5	2.0	2.0	2.5	NaN	NaN	35.0
27	2.0	3.0	3.0	NaN	none	33.0
26	2.0	4.5	4.5	NaN	tcf	NaN
34	3.0	2.0	2.5	2.1	tc	40.0
30	3.0	4.5	4.5	5.0	none	40.0
31	3.0	3.0	2.0	2.5	tc	40.0
7	3.0	6.9	4.8	2.3	NaN	40.0
4	3.0	4.5	4.5	5.0	NaN	40.0
25	3.0	2.0	2.0	2.0	none	40.0
38	3.0	2.0	2.5	2.0	NaN	37.0
13	3.0	3.5	4.0	5.1	tcf	37.0
10	3.0	3.5	4.0	4.6	none	36.0
29	3.0	2.0	2.5	NaN	NaN	35.0
22	3.0	3.5	4.0	4.6	tcf	27.0
3	3.0	3.7	4.0	5.0	tc	NaN
6	3.0	4.0	5.0	5.0	tc	NaN
2	NaN	NaN	NaN	NaN	NaN	38.0

1.5 DF Selecting

Selecting is similar to the slicing. Typically, it deals with the following: 1. We are typically interested in extracting rows based on conditions. 2. Conditions are based on the values of one or more columns.

Consider the dataset given in Basic-1.csv. Display all rows related to White-collar jobs.

```
[26]: import pandas as pd
df = pd.read_csv('data/Basic-1.csv', delimiter=',')
display(df.head())

seleted_rows = df['Job Type']=='White-collar'
display(df.loc[seleted_rows,:])
```

	Gender	Job Type	Province
0	M	Pink-collar	Hejaz
1	M	White-collar	Central
2	M	Pink-collar	Hejaz
3	M	Pink-collar	Hejaz
4	M	Gold-collar	Eastern

	Gender	Job Type	Province
1	M	White-collar	Central
5	M	White-collar	Hejaz
8	M	White-collar	Eastern
9	F	White-collar	Hejaz
10	M	White-collar	Hejaz

```

...      ...      ...      ...
1266      M  White-collar  Hejaz
1267      M  White-collar  Hejaz
1271      F  White-collar  Central
1272      M  White-collar  Central
1276      F  White-collar  Central

```

[568 rows x 3 columns]

Consider the dataset given in Basic-1.csv. 1. Display all rows related to White-collar jobs in Eastern Province. 2. Display the statistical summary of Gender column in the above selection.

```

[27]: import pandas as pd
df = pd.read_csv('data/Basic-1.csv', delimiter=',')

seleted_rows = (df['Province']=='Eastern') & (df['Job Type']=='White-collar')
display(df.loc[seleted_rows,:])

# Summary
display(df.loc[seleted_rows,:].describe().iloc[:,0])
# display(df.loc[seleted_rows,:].describe())
# display(df.loc[seleted_rows,:].describe().loc[:, 'Gender'])

```

```

      Gender      Job Type Province
8          M  White-collar  Eastern
18         F  White-collar  Eastern
27         M  White-collar  Eastern
31         M  White-collar  Eastern
37         M  White-collar  Eastern
...      ...      ...      ...
1226      M  White-collar  Eastern
1232      M  White-collar  Eastern
1238      F  White-collar  Eastern
1241      M  White-collar  Eastern
1256      M  White-collar  Eastern

```

[138 rows x 3 columns]

```

count      138
unique       2
top         M
freq       98
Name: Gender, dtype: object

```

1.5.1 Example: Displaying Data

Question-A: Consider the dataset given in Basic-1.csv. 1. Display all rows related to White-collar and Blue-collar jobs. 2. Display all rows related to White-collar and Blue-collar jobs for Females.

3. Display the statistical summary of Province column in the above selection. 4. Among all the White-collar and Blue-collar job Females, what proportion works in Eastern Province.

[28]: *#1. Display all rows related to White-collar and Blue-collar jobs.*

```
import pandas as pd
df = pd.read_csv('data/Basic-1.csv', delimiter=',')
display(df)

seleted_rows = df['Job Type'].isin(['White-collar','Blue-collar'])

# seleted_rows = (df['Job Type']=='White-collar') / (df['Job
↪Type']=='Blue-collar')

display(df.loc[seleted_rows,:])
```

	Gender	Job Type	Province
0	M	Pink-collar	Hejaz
1	M	White-collar	Central
2	M	Pink-collar	Hejaz
3	M	Pink-collar	Hejaz
4	M	Gold-collar	Eastern
...
1273	M	Blue-collar	Central
1274	M	Pink-collar	Hejaz
1275	F	Blue-collar	Hejaz
1276	F	White-collar	Central
1277	F	Blue-collar	Eastern

[1278 rows x 3 columns]

	Gender	Job Type	Province
1	M	White-collar	Central
5	M	White-collar	Hejaz
6	M	Blue-collar	Eastern
7	M	Blue-collar	Central
8	M	White-collar	Eastern
...
1272	M	White-collar	Central
1273	M	Blue-collar	Central
1275	F	Blue-collar	Hejaz
1276	F	White-collar	Central
1277	F	Blue-collar	Eastern

[984 rows x 3 columns]

[29]: *# 2. Display all rows related to White-collar and Blue-collar jobs for Females.*

```

import pandas as pd
df = pd.read_csv('data/Basic-1.csv', delimiter=',')

seleted_rows = df['Job Type'].isin(['White-collar', 'Blue-collar']) & 
↳ (df['Gender']=='F')

display(df.loc[seleted_rows,:])

# 3. Display the statistical summary of Province column in the above selection.
df.loc[seleted_rows,['Province']].describe()

```

	Gender	Job Type	Province
9	F	White-collar	Hejaz
11	F	Blue-collar	Central
13	F	White-collar	Central
14	F	White-collar	Hejaz
15	F	White-collar	Central
...
1254	F	White-collar	Hejaz
1271	F	White-collar	Central
1275	F	Blue-collar	Hejaz
1276	F	White-collar	Central
1277	F	Blue-collar	Eastern

[311 rows x 3 columns]

```

[29]:      Province
count      311
unique        3
top      Central
freq       136

```

```

[30]: # 4. Among all the White-collar and Blue-collar job Females, what proportion
↳ works in Eastern Province.

ndf = df.loc[seleted_rows,'Province']
display(ndf)
print(ndf.value_counts())

proportion = ndf.value_counts()[2]/len(ndf.index)*100

## to get length of dataframe
# print(len(ndf.index))
# print(ndf.shape[0])
# print(len(ndf))

```

```
print(f'The proportion of Females working as blue/white collar in Eastern_
↪ province is: {proportion: 0.2f}%')
```

```
9      Hejaz
11     Central
13     Central
14     Hejaz
15     Central
...
1254   Hejaz
1271   Central
1275   Hejaz
1276   Central
1277   Eastern
Name: Province, Length: 311, dtype: object
```

```
Province
Central    136
Hejaz      104
Eastern     71
Name: count, dtype: int64
The proportion of Females working as blue/white collar in Eastern province is:
22.83%
```

```
/tmp/ipykernel_21727/2696248615.py:6: FutureWarning: Series.__getitem__ treating
keys as positions is deprecated. In a future version, integer keys will always
be treated as labels (consistent with DataFrame behavior). To access a value by
position, use `ser.iloc[pos]`
    proportion = ndf.value_counts()[2]/len(ndf.index)*100
```

1.6 Addition DF Methods

1.6.1 Count Rows and Columns

Consider the dataset given in Basic-2.csv. 1. Count the number of rows and columns. 2. Count the number of non-null rows for each column.

```
[31]: import pandas as pd
df = pd.read_excel('data/Basic-2.xlsx')
display(df.head())

print(f'The number of rows are {len(df.index)}, and the number of columns are_
↪ {len(df.columns)}')

print(f'The number of non-null rows for each column are:\n{df.count()}')

print(f'The number of null rows for each column are:\n{df.isna().sum()}')

print(df.shape)
```


	c1	c2	c3	c4	c5	c6		c7	c8	c9	c10	c11	\
0	1.0	5.0	NaN	NaN	NaN	40.0		NaN	NaN	2.0	NaN	11.0	
1	2.0	4.5	5.8	NaN	NaN	35.0	ret_allw	NaN	NaN	yes	11.0		
2	NaN	NaN	NaN	NaN	NaN	38.0	empl_contr	NaN	5.0	NaN	11.0		
3	3.0	3.7	4.0	5.0	tc	NaN		NaN	NaN	NaN	yes	NaN	
4	3.0	4.5	4.5	5.0	NaN	40.0		NaN	NaN	NaN	NaN	12.0	

		c12	c13	c14	c15	c16	c17
0		average	NaN	NaN	yes	NaN	good
1	below	average	NaN	full	NaN	full	good
2		generous	yes	half	yes	half	good
3		NaN	NaN	NaN	yes	NaN	good
4		average	NaN	half	yes	half	good

The number of rows are 40, and the number of columns are 17

The number of non-null rows for each column are:

c1	39
c2	39
c3	30
c4	12
c5	24
c6	37
c7	18
c8	7
c9	24
c10	18
c11	38
c12	37
c13	16
c14	25
c15	20
c16	24
c17	40

dtype: int64

The number of null rows for each column are:

c1	1
c2	1
c3	10
c4	28
c5	16
c6	3
c7	22
c8	33
c9	16
c10	22
c11	2
c12	3
c13	24

```

c14      15
c15      20
c16      16
c17       0
dtype: int64
(40, 17)

```

1.6.2 Apply Lambda Functions

Lambda functions can be applied to each row or column using `.apply()` method.

Consider the data in file Basic-2-Clean.csv. 1. In columns c1, c6 and c11, convert every number to the nearest integer. 2. In column c12, replace the space between two words by underscore.

```

[32]: # 1. In columns c1, c6 and c11, convert every number to the nearest integer.
import pandas as pd
df = pd.read_csv('data/Basic-2-Clean.csv')
display(df.head())

df['c1']=df['c1'].apply(lambda x: int(round(x))) # apply for single column
# df['c1']= (lambda x: round(x))(df['c1']) #another style

df.loc[:,['c6','c11']]=df.loc[:,['c6','c11']].applymap(lambda x: int(round(x)))
↪ #for multiple columns

display(df.head())

```

	c1	c2	c3	c5	c6	c9	c11 \
0	1.000000	5.000000	3.913333	none	40.000000	2.000000	11.000000
1	2.000000	4.500000	5.800000	none	35.000000	4.583333	11.000000
2	2.102564	3.620513	3.913333	none	38.000000	5.000000	11.000000
3	3.000000	3.700000	4.000000	tc	37.810811	4.583333	11.105263
4	3.000000	4.500000	4.500000	none	40.000000	4.583333	12.000000

	c12	c14	c16	c17
0	average	half	full	good
1	below average	full	full	good
2	generous	half	half	good
3	below average	half	full	good
4	average	half	half	good

/tmp/ipykernel_21727/3215444923.py:9: FutureWarning: DataFrame.applymap has been deprecated. Use DataFrame.map instead.

```

df.loc[:,['c6','c11']]=df.loc[:,['c6','c11']].applymap(lambda x: int(round(x)))
) #for multiple columns

```

	c1	c2	c3	c5	c6	c9	c11	c12	c14 \
0	1	5.000000	3.913333	none	40.0	2.000000	11.0	average	half
1	2	4.500000	5.800000	none	35.0	4.583333	11.0	below average	full
2	2	3.620513	3.913333	none	38.0	5.000000	11.0	generous	half

3	3	3.700000	4.000000	tc	38.0	4.583333	11.0	below average	half
4	3	4.500000	4.500000	none	40.0	4.583333	12.0	average	half

	c16	c17
0	full	good
1	full	good
2	half	good
3	full	good
4	half	good

[33]: # 2. In column c12, replace the space between two words by underscore.

```
display(df.head())
df['c12']=df['c12'].apply(lambda x: x.replace(' ','_'))
display(df.head())
```

	c1	c2	c3	c5	c6	c9	c11	c12	c14	\
0	1	5.000000	3.913333	none	40.0	2.000000	11.0	average	half	
1	2	4.500000	5.800000	none	35.0	4.583333	11.0	below average	full	
2	2	3.620513	3.913333	none	38.0	5.000000	11.0	generous	half	
3	3	3.700000	4.000000	tc	38.0	4.583333	11.0	below average	half	
4	3	4.500000	4.500000	none	40.0	4.583333	12.0	average	half	

	c16	c17
0	full	good
1	full	good
2	half	good
3	full	good
4	half	good

	c1	c2	c3	c5	c6	c9	c11	c12	c14	\
0	1	5.000000	3.913333	none	40.0	2.000000	11.0	average	half	
1	2	4.500000	5.800000	none	35.0	4.583333	11.0	below_average	full	
2	2	3.620513	3.913333	none	38.0	5.000000	11.0	generous	half	
3	3	3.700000	4.000000	tc	38.0	4.583333	11.0	below_average	half	
4	3	4.500000	4.500000	none	40.0	4.583333	12.0	average	half	

	c16	c17
0	full	good
1	full	good
2	half	good
3	full	good
4	half	good

1.6.3 Apply General Functions

Similar to Lambda functions, user defined or inbuilt functions can be applied to each row or column using `.apply()` method. Moreover, `.applymap()` method can be used for multiple columns.

Consider the data in file Basic-2-Clean.csv.

In columns c2, c3 and c9, round the values to the nearest 0, 0.5 or 1. If any value is negative, then replace it with zero.

```
[34]: def custom_round(x):
    if x <=0:
        return 0
    int_x = int(x)
    if (x <= int_x+0.25):
        return int_x
    elif (x > int_x+0.25) and (x <= int_x+0.75):
        return int_x+0.5
    else:
        return int_x+1

import pandas as pd
df = pd.read_csv('data/Basic-2-Clean.csv')
display(df.head())
# df['c2']=df['c2'].apply(custom_round) # for single column
df.loc[:,['c2','c3','c9']]=df.loc[:,['c2','c3','c9']].applymap(custom_round) #
# for multiple columns
display(df.head())
```

	c1	c2	c3	c5	c6	c9	c11 \
0	1.000000	5.000000	3.913333	none	40.000000	2.000000	11.000000
1	2.000000	4.500000	5.800000	none	35.000000	4.583333	11.000000
2	2.102564	3.620513	3.913333	none	38.000000	5.000000	11.000000
3	3.000000	3.700000	4.000000	tc	37.810811	4.583333	11.105263
4	3.000000	4.500000	4.500000	none	40.000000	4.583333	12.000000

	c12	c14	c16	c17
0	average	half	full	good
1	below average	full	full	good
2	generous	half	half	good
3	below average	half	full	good
4	average	half	half	good

/tmp/ipykernel_21727/1219133502.py:16: FutureWarning: DataFrame.applymap has been deprecated. Use DataFrame.map instead.

```
df.loc[:,['c2','c3','c9']]=df.loc[:,['c2','c3','c9']].applymap(custom_round)
# for multiple columns
```

	c1	c2	c3	c5	c6	c9	c11	c12	c14 \
0	1.000000	5.0	4.0	none	40.000000	2.0	11.000000	average	half
1	2.000000	4.5	6.0	none	35.000000	4.5	11.000000	below average	full
2	2.102564	3.5	4.0	none	38.000000	5.0	11.000000	generous	half
3	3.000000	3.5	4.0	tc	37.810811	4.5	11.105263	below average	half
4	3.000000	4.5	4.5	none	40.000000	4.5	12.000000	average	half

	c16	c17
0	full	good
1	full	good
2	half	good
3	full	good
4	half	good

2 Data Visualization

Graphical representation of data

2.1 Why Visualization

1. When data is small (no magic numbers, but say 5 rows and 2 columns), then one may extract meaningful patterns by looking at the data.
2. However, when we have large data (again no magic numbers, but say 5000 rows and 2000 columns), then mere observation may not be fruitful in identifying patterns.
3. Visualization is the first typical step in data analysis.

2.2 Typical Visualizations

1. Histogram
2. Box Plots
3. Scatter Plots
4. Line Plots

2.3 Pandas Plots

1. Pandas provide in-built basic plots using the numpy and matplotlib libraries.
2. Several plots from pandas are under `df.plot`, which can be accessed via *kind* option.
3. In addition to that, specific plots like histogram (`df.hist()`), and boxplot (`df.boxplot()`) are also available.
4. However, pandas basic plots are not enough, and in the course we look at the seaborn library for plotting.

In the following cells, we will look at most typical plots that can be constructed using seaborn plots.

2.4 Seaborn Plots

1. Seaborn library provides a variety of plots, including histogram, boxplots, lineplots, scatterplots, countplot, violinplot, swarmplot, pairplot, catplot etc.
2. These plots can be accessed through *.countplot*, *.violinplot*, *catplot*, etc.
3. Typical parameters include

x, y, and/or hue

where *x*, *y* refers to the x and y axis, and *hue* defines subsets of the data, which will be drawn on separate facets in the grid.

4. It also comes with huge customization.

5. The Seaborn Plots are aesthetically better than basic plots from pandas library.
6. The library requires numpy and matplotlib libraries.
7. It is integrated with pandas data.

In the following cell, we will look at most typical plots from the seaborn library.

2.4.1 Example: Seaborn's Basic Plots

Question-B: Consider the data in file Basic-3.csv.

1. Read and display information the data.
2. Add a new column to dataframe that represents total score. The total score is sum of score of math, reading, and writing scores.
3. Draw histograms of both numeric and non-numeric columns.
4. Draw box-plots for the numeric columns, and differentiate by *test preparation course* column.
5. Draw three overlapping scatter plots of columns *math score*, *reading score* and *writing score* w.r.t *Total score*. Set the markers transparency level (alpha) to 0.5.
6. Draw scatter plot of columns *math score* and *writing score*, where the size of the marker is based on column *Total score*. In addition to that, reduce the marker transparency (alpha) to 0.5.
7. Plot *Total score* in ascending order on the x axis, and the corresponding pairwise absolute differences of *math score*, *reading score* and *writing score* on the y axis.

[35]: # 1. Read and display information the data.

```
import pandas as pd
df = pd.read_csv('data/Basic-3.csv')

display(df.head())
display(df.info())
```

	gender	race/ethnicity	parental level of education	lunch	\
0	female	group B	bachelor's degree	standard	
1	female	group C	some college	standard	
2	female	group B	master's degree	standard	
3	male	group A	associate's degree	free/reduced	
4	male	group C	some college	standard	

	test preparation course	math score	reading score	writing score
0	none	72	72	74
1	completed	69	90	88
2	none	90	95	93
3	none	47	57	44
4	none	76	78	75

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1000 entries, 0 to 999
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----

```

0    gender                1000 non-null    object
1    race/ethnicity        1000 non-null    object
2    parental level of education 1000 non-null    object
3    lunch                 1000 non-null    object
4    test preparation course 1000 non-null    object
5    math score            1000 non-null    int64
6    reading score         1000 non-null    int64
7    writing score          1000 non-null    int64
dtypes: int64(3), object(5)
memory usage: 62.6+ KB

```

None

```

[36]: # 2. Add a new column to dataframe that represents total score.
      # The total score is sum of score of math, reading, and writing scores

      df['Total score']=df['math score']+df['reading score']+df['writing score']

      df

```

```

[36]:      gender race/ethnicity parental level of education      lunch \
0    female      group B      bachelor's degree      standard
1    female      group C      some college      standard
2    female      group B      master's degree      standard
3     male      group A      associate's degree  free/reduced
4     male      group C      some college      standard
..     ...      ...      ...      ...
995  female      group E      master's degree      standard
996   male      group C      high school  free/reduced
997  female      group C      high school  free/reduced
998  female      group D      some college      standard
999  female      group D      some college  free/reduced

      test preparation course  math score  reading score  writing score \
0                none          72          72          74
1          completed          69          90          88
2                none          90          95          93
3                none          47          57          44
4                none          76          78          75
..                ...      ...      ...      ...
995          completed          88          99          95
996                none          62          55          55
997          completed          59          71          65
998          completed          68          78          77
999                none          77          86          86

      Total score
0              218

```

```

1          247
2          278
3          148
4          229
..         ...
995        282
996        172
997        195
998        223
999        249

```

```
[1000 rows x 9 columns]
```

```
[37]: %pip install seaborn
```

```

Collecting seaborn
  Downloading seaborn-0.13.2-py3-none-any.whl.metadata (5.4 kB)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in
/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (from
seaborn) (2.2.3)
Requirement already satisfied: pandas>=1.2 in
/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (from
seaborn) (2.2.3)
Collecting matplotlib!=3.6.1,>=3.4 (from seaborn)
  Downloading matplotlib-3.10.1-cp311-cp311-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (11 kB)
Collecting contourpy>=1.0.1 (from matplotlib!=3.6.1,>=3.4->seaborn)
  Downloading contourpy-1.3.1-cp311-cp311-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (5.4 kB)
Collecting cycler>=0.10 (from matplotlib!=3.6.1,>=3.4->seaborn)
  Downloading cycler-0.12.1-py3-none-any.whl.metadata (3.8 kB)
Collecting fonttools>=4.22.0 (from matplotlib!=3.6.1,>=3.4->seaborn)
  Downloading fonttools-4.56.0-cp311-cp311-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (101 kB)
                                101.9/101.9 kB
682.0 kB/s eta 0:00:00a 0:00:01
Collecting kiwisolver>=1.3.1 (from matplotlib!=3.6.1,>=3.4->seaborn)
  Downloading kiwisolver-1.4.8-cp311-cp311-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (6.2 kB)
Requirement already satisfied: packaging>=20.0 in
/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn) (24.2)
Collecting pillow>=8 (from matplotlib!=3.6.1,>=3.4->seaborn)
  Downloading pillow-11.1.0-cp311-cp311-manylinux_2_28_x86_64.whl.metadata (9.1
kB)
Collecting pyparsing>=2.3.1 (from matplotlib!=3.6.1,>=3.4->seaborn)
  Downloading pyparsing-3.2.1-py3-none-any.whl.metadata (5.0 kB)
Requirement already satisfied: python-dateutil>=2.7 in

```



```

/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (from
matplotlib!=3.6.1,>=3.4->seaborn) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in
/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (from
pandas>=1.2->seaborn) (2025.1)
Requirement already satisfied: tzdata>=2022.7 in
/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (from
pandas>=1.2->seaborn) (2025.1)
Requirement already satisfied: six>=1.5 in
/home/motid/kfupm/assignment2-ICS/.venv/lib/python3.11/site-packages (from
python-dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn) (1.17.0)
Downloading seaborn-0.13.2-py3-none-any.whl (294 kB)
294.9/294.9 kB
2.1 MB/s eta 0:00:00a 0:00:01
Downloading
matplotlib-3.10.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl
(8.6 MB)
8.6/8.6 MB
10.4 MB/s eta 0:00:0000:0100:01
Downloading
contourpy-1.3.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (326
kB)
326.2/326.2 kB
9.5 MB/s eta 0:00:00ta 0:00:01
Downloading cycler-0.12.1-py3-none-any.whl (8.3 kB)
Downloading
fonttools-4.56.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (4.9
MB)
4.9/4.9 MB
13.1 MB/s eta 0:00:0000:0100:01
Downloading
kiwisolver-1.4.8-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1.4
MB)
1.4/1.4 MB
12.2 MB/s eta 0:00:0000:0100:01
Downloading pillow-11.1.0-cp311-cp311-manylinux_2_28_x86_64.whl (4.5 MB)
4.5/4.5 MB
13.7 MB/s eta 0:00:0000:0100:01
Downloading pyparsing-3.2.1-py3-none-any.whl (107 kB)
107.7/107.7 kB
5.4 MB/s eta 0:00:00
Installing collected packages: pyparsing, pillow, kiwisolver, fonttools,
cycler, contourpy, matplotlib, seaborn
Successfully installed contourpy-1.3.1 cycler-0.12.1 fonttools-4.56.0
kiwisolver-1.4.8 matplotlib-3.10.1 pillow-11.1.0 pyparsing-3.2.1 seaborn-0.13.2

[notice] A new release of pip is
available: 24.0 -> 25.0.1

```

[notice] To update, run:

```
pip install --upgrade pip
```

Note: you may need to restart the kernel to use updated packages.


```
[38]: %matplotlib inline
      # 3. (a) Draw histograms of both numeric and non-numeric columns.

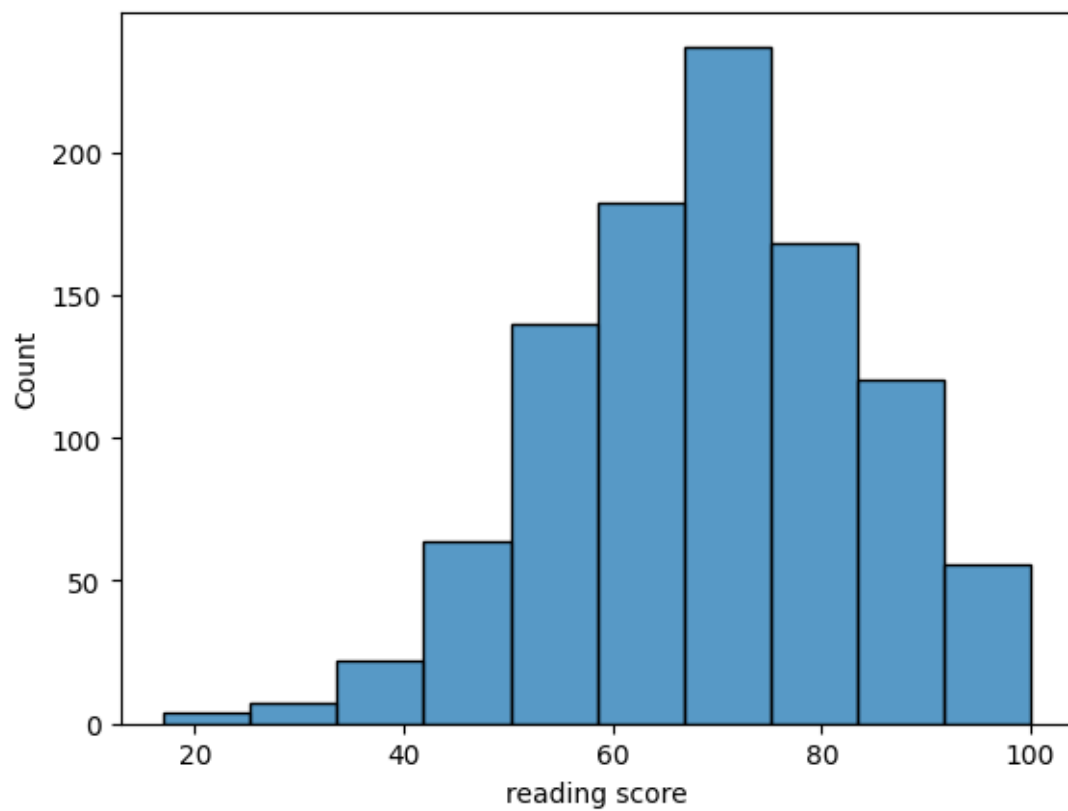
      # Load all the libraries
      import matplotlib.pyplot as plt
      import seaborn as sns


      # Identify numeric columns
      num_columns = df.select_dtypes(exclude='object').columns


      for c in num_columns:
          plt.figure()
          sns.histplot(x=c, bins=10, data=df);
          plt.show()

      # sns.histplot(x=df["column_name"], bins=15, kde=True, color='red',
      ↪stat='density')
      # plt.show()
      ## To draw a single column
      # plt.figure()
      # sns.histplot(x='math score', bins=10, data=df);
      # plt.show()
```










```
[39]: # pip install seaborn
```

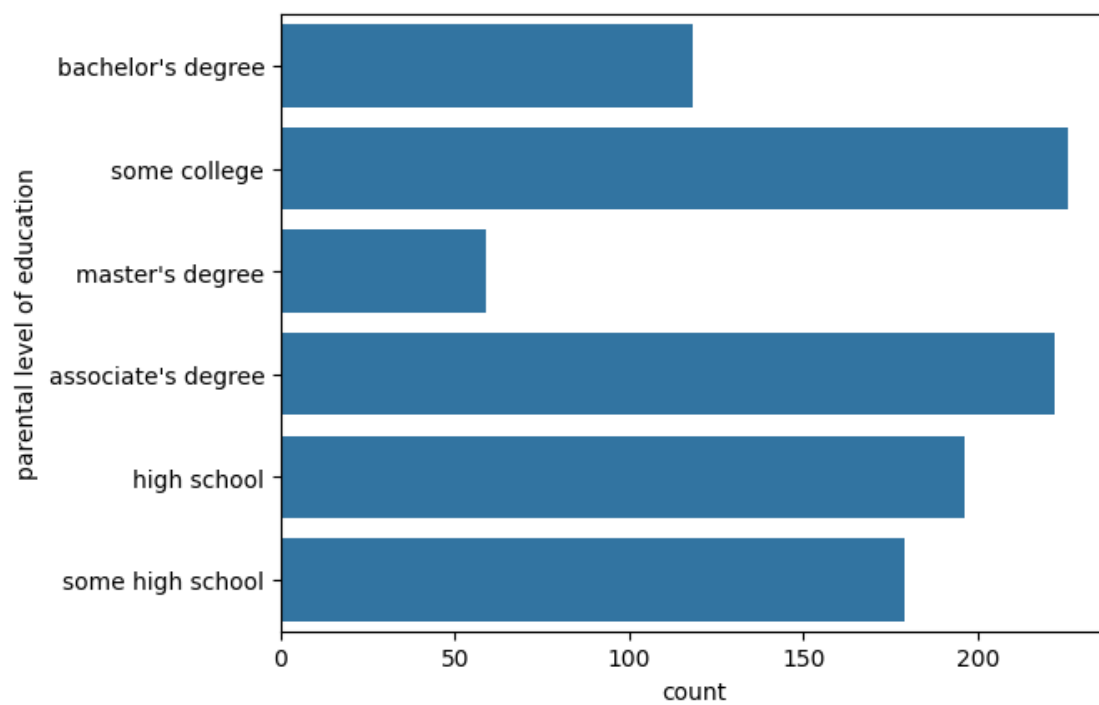

```
[40]: # (b) For non-numeric columns, count plot option may be used.
```



```
# Identify non-numeric columns
obj_columns = df.select_dtypes('object').columns

# Plot for each identified column
for c in obj_columns:
    plt.figure()
    #     sns.histplot(y=c,data=df);
    sns.countplot(y=c,data=df);
    plt.show()

## To draw a single column
# plt.figure()
# sns.countplot(y='gender',data=df);
# plt.show()
```







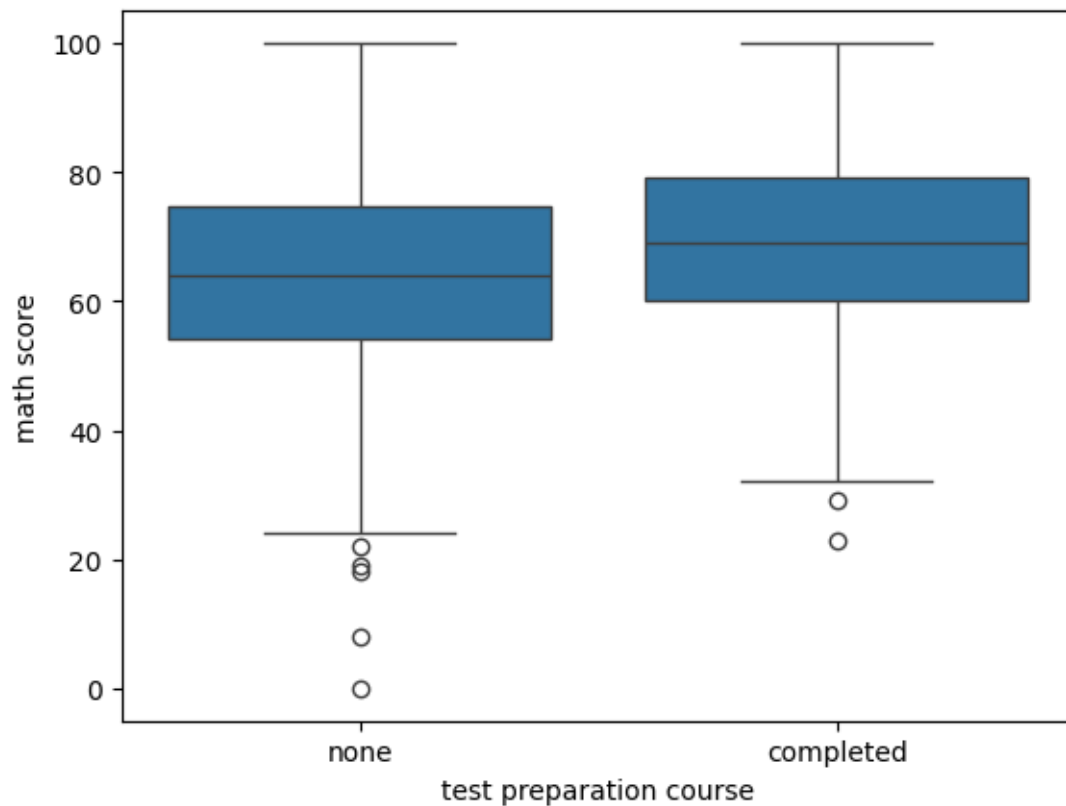
```
[41]: # 4. Draw box-plots for the numeric columns, and differentiate by *test_
      ↪ preparation course* column.
```


```
import matplotlib.pyplot as plt
import seaborn as sns


# Identify numeric columns
num_columns = df.select_dtypes(exclude='object').columns


for c in num_columns:
    plt.figure()
    sns.boxplot(y=c,x='test preparation course',data=df);
    plt.show()

# plt.figure()
# sns.boxplot(y='math score',x='test preparation course',data=df);
# plt.show()
```



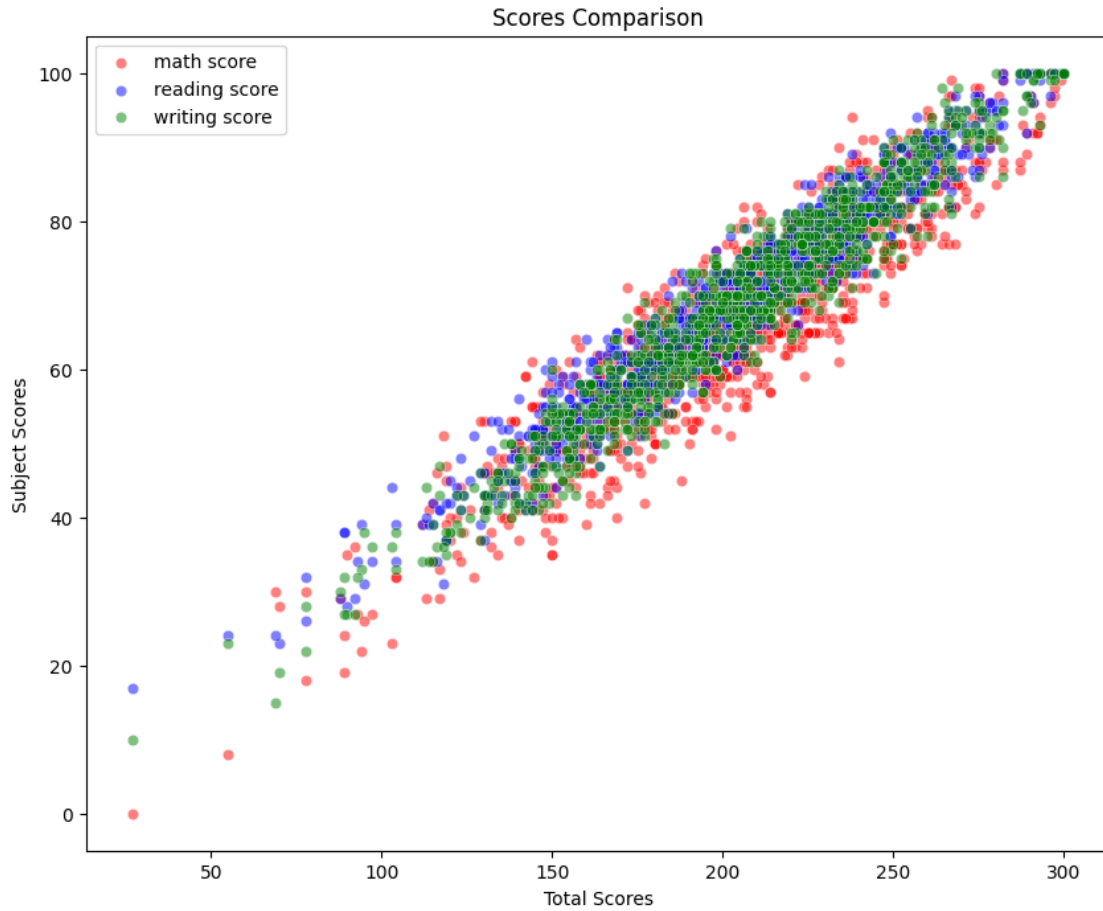







```
[42]: # 5. Draw three overlapping scatter plots of columns *math score, reading_
      ↪score* and *writing score*
      #w.r.t *Total score*. Set the markers transparency level (alpha) to 0.5.

      plt.figure(figsize=(10,8))
      sns.scatterplot(x='Total score', y='math score', color='red', label='math_
      ↪score', alpha=0.5, data=df)
      sns.scatterplot(x='Total score', y='reading score',color='blue',label='reading_
      ↪score', alpha=0.5, data=df)
      sns.scatterplot(x='Total score', y='writing score',color='green',label='writing_
      ↪score', alpha=0.5, data=df)
      plt.ylabel("Subject Scores");
      plt.xlabel("Total Scores");
      plt.title("Scores Comparison");
      plt.show()
```



```
[43]: # 6. Draw scatter plot of columns *math score* and *writing score*,  
#where the size of the marker is based on column *Total score*.  
#In addition to that, reduce the marker transparency (alpha) to 0.5.
```

```
plt.figure(figsize=(10,10))  
sns.scatterplot(x='math score', y='writing score',  
               size='Total score', sizes=(20,200),  
               alpha=0.5, color='red',  
               data=df)  
plt.show()
```



2.4.2 Example: Seaborn's Advanced Plots

Question-C: Consider the data in file Basic-4-Clean.csv.

1. Read and display information the data. Remove all rows containing NA values.
2. Depict *Category* column by count. In another plot, depict *Category* counts (differentiated) by *Type*.
3. Draw a boxplot plot of *Installs* vs *Rating* columns, ordered by *Installs*.
4. Draw a violinplot plot of *Content Rating* vs *Rating* columns, differentiated w.r.t *Type* .
5. Plot a barplot depicting count of *Category* that have *Installs* above 0.5B.
6. Depicting the *Size* of *Category* that have *Installs* above 0.5B.
7. For all the paid apps, depict *Genres* by count differentiated by *Content Rating*. In another plot depict *Genres* by *Rating*.

```
[44]: # 1. Read and display information the data. Remove all rows containing NA
      ↪ values.
```

```
import pandas as pd
df = pd.read_csv('data/Basic-4-Clean.csv')
print(df.info())
#print(df.shape)# Can also be obtained from info
#print(df.count())# Can also be obtained from info
display(df.sample(5))
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7723 entries, 0 to 7722
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   App                    7723 non-null  object
1   Category               7723 non-null  object
2   Rating                 7723 non-null  float64
3   Reviews                7723 non-null  int64
4   Size                   7723 non-null  float64
5   Installs               7723 non-null  object
6   Type                   7723 non-null  object
7   Price                  7723 non-null  float64
8   Content Rating         7723 non-null  object
9   Genres                 7723 non-null  object
10  Last Updated           7723 non-null  object
11  Current Ver            7723 non-null  object
12  Android Ver            7723 non-null  object
dtypes: float64(3), int64(1), object(9)
memory usage: 784.5+ KB
None
```

	App	Category	Rating	Reviews	Size	\
322	OkCupid Dating	DATING	4.1	285726	15000.0	
1446	Learn To Draw Glow Flower	FAMILY	4.4	7320	10000.0	
5229	CJ WOW SHOP	SHOPPING	4.2	2099	4.0	
7409	FG Mobile	FAMILY	3.3	130	13000.0	
6077	Download Manager - File & Video	TOOLS	3.9	8780	5.0	

	Installs	Type	Price	Content Rating	Genres	\
322	10000000+	Free	0.0	Mature 17+	Dating	
1446	1000000+	Free	0.0	Everyone	Entertainment;Education	
5229	100000+	Free	0.0	Everyone	Shopping	
7409	5000+	Free	0.0	Everyone	Education	
6077	1000000+	Free	0.0	Everyone	Tools	


	Last Updated	Current Ver	Android Ver
322	30-Jul-18	11.10.1	4.1 and up


1446	7-Jul-17	1.0.1	4.0.3 and up
5229	23-Dec-17	1.1.1	4.0.3 and up
7409	21-Dec-17	2.0.0	4.0 and up
6077	13-Jun-18	2.7.5	4.2 and up

[45]: # 2. Depict *Category* column by count. In another plot, depict *Category* ↪ counts (differentiated) by *Type*.

```
#plot category counts
import seaborn as sns
plt.figure(figsize=(10,10))
ax = sns.countplot(y="Category", data=df)
plt.title('Apps available')
plt.show()

#plot category with type
import seaborn as sns
plt.figure(figsize=(10,10))
ax = sns.countplot(y='Category', hue='Type', data=df)
plt.title('Apps popularity')
plt.show()
```

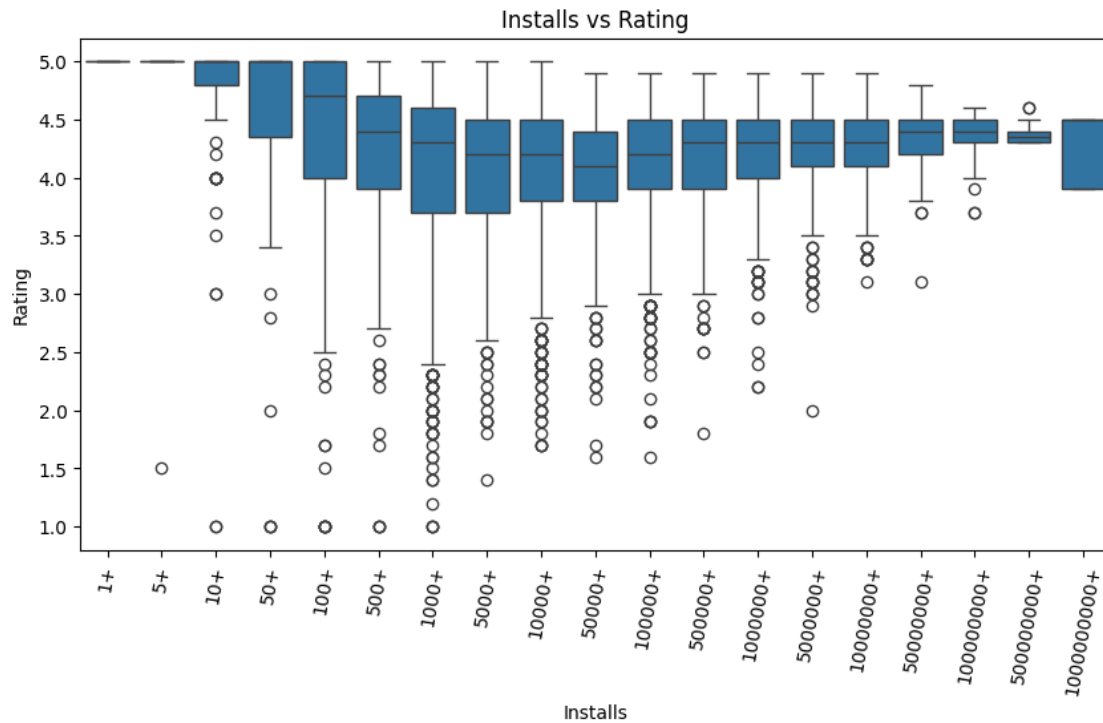




```
[46]: #3. Draw a boxplot plot of *Installs* vs *Rating* columns, ordered by
      ↳*Installs*.

      #adding a new column
      df['Installs_num']=df['Installs'].apply(lambda x:x.replace('+','')).apply(pd.
      ↳to_numeric)


      #the box plot
      ax = plt.figure(figsize=(10,5))
      sns.boxplot(x="Installs", y="Rating", data=df.sort_values(by=['Installs_num']))
      #,palette="Set2"
      plt.title("Installs vs Rating")
      plt.xticks(rotation=80)
      plt.show()
```



```
[47]: #4. Draw a violinplot plot of *Content Rating* vs *Rating* columns,
      ↪ differentiated w.r.t *Type* .

import seaborn as sns
plt.figure(figsize=(10,10))
sns.violinplot(x='Content Rating', y='Rating', hue='Type', data=df)
plt.title('Apps ratings')
plt.show()
#violinplot plot also depicts the distribution.

# # contrast with box plots
# plt.figure(figsize=(10,10))
# sns.boxplot(x='Content Rating', y='Rating', hue='Type', data=df)
# plt.title('Apps ratings')
# plt.show()
```




```
[48]: #5. Plot a barplot depicting count of *Category* that have *Installs* above 0.5B.
import seaborn as sns
seleted_rows = (df['Installs']=='1000000000+') | (df['Installs']=='500000000+')

plt.figure(figsize=(10,2))
sns.countplot(y='Category', data = df.loc[seleted_rows,:],palette="Set2")
plt.title('App Categories in 0.5B+ & above')
plt.show()
```

/tmp/ipykernel_21727/394882582.py:6: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(y='Category', data = df.loc[seleted_rows,:],palette="Set2")
```



[49]: #6. Depicting the *Size* of *Category* that have *Installs* above 0.5B.


```
import seaborn as sns

plt.figure(figsize=(5,5))
# ax=sns.swarmplot(x='Category', y='Size', data = df.loc[seleted_rows,:])
# ax=sns.violinplot(x='Category', y='Size', data = df.loc[seleted_rows,:])
ax=sns.boxplot(x='Category', y='Size', data = df.loc[seleted_rows,:
↵],palette="Set2")
plt.xticks(rotation=80)
plt.title('Top Apps sizes in Kb')
plt.show()
```

/tmp/ipykernel_21727/3301675080.py:8: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
ax=sns.boxplot(x='Category', y='Size', data =
df.loc[seleted_rows,:],palette="Set2")
```



3 References:

Main ref. ISE 291

3.1 Data Sets:

1. Basic-2: Labor Relations Data, UCI Machine Learning Repository [<http://archive.ics.uci.edu/ml>]. Irvine, CA: University of California, School of Information and Computer Science. Published in: Bergadano, F., Matwin, S., Michalski, R., Zhang, J., Measuring Quality of Concept Descriptions, Procs. of the 3rd European Working Sessions on Learning, Glasgow, October 1988.

2. Basic-3: <https://www.kaggle.com/spscientist/students-performance-in-exams>
3. Basic-4: *modified*, <https://www.kaggle.com/lava18/google-play-store-apps>

3.2 Others:

1. Series: <https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.Series.html>
2. DataFrames: <https://pandas.pydata.org/pandas-docs/stable/reference/frame.html>
3. Read CSV: https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.read_csv.html
4. Read XLSX: https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.read_excel.html
5. Visualization: https://pandas.pydata.org/pandas-docs/stable/user_guide/visualization.html
6. Seaborn: <https://seaborn.pydata.org/>